

STAT 423 Winter 2026 - Final Project

Krish Doshi, Lauren Yee, Kian Sharifzadeh

2026-03-09

Data Processing - Krish Doshi

```
orig_data <- read.csv("billboard.csv")

cleaned_data <- data.frame(artist = orig_data$artist, rating = orig_data$overall_rating,
  age=orig_data$front_person_age)

cleaned_data$gender <- ifelse(orig_data$artist_male == 1, "Male", "Not Male")

cleaned_data$race <- ifelse(orig_data$artist_white == 1, "White", "Not White")

cleaned_data$genre <- ifelse(is.na(orig_data$cdr_genre), "Genre Not Specified",
  orig_data$cdr_genre)

cleaned_data$genre <- ifelse(grepl("Pop", orig_data$cdr_genre), "Pop",
  ifelse(grepl("Country", orig_data$cdr_genre), "Country",
  ifelse(grepl("Hip-Hop", orig_data$cdr_genre), "Hip-Hop",
  ifelse(grepl("Rock", orig_data$cdr_genre), "Rock",
  ifelse(grepl("Jazz", orig_data$cdr_genre), "Jazz", "Other")))))

cleaned_data <- cbind(cleaned_data, bpm=orig_data$bpm, energy=orig_data$energy,
  danceability = orig_data$danceability, happiness = orig_data$happiness,
  loudness_db = orig_data$loudness_d_b,
  length_sec = orig_data$length_sec)

cleaned_data$year <- as.numeric(substr(orig_data$date, 1,
  unlist(gregexpr("-", orig_data$date))[1] - 1))

cleaned_data <- cleaned_data[rowSums(is.na(cleaned_data)) == 0 , ]

model_results <- list()

for(i in seq(1960, 2030, 10)){
  decade_label <- paste0(i - 10, "s")

  subpop <- cleaned_data[cleaned_data$year >= i - 10 &
    cleaned_data$year < i, ]

  # --- SAFETY CHECK 1: Ensure enough rows exist ---
  # If you have < 50 rows, the interaction model will likely fail.
```

```

if(nrow(subpop) < 50) {
  message(paste("Skipping", decade_label, "- Not enough data points."))
  next
}

factor_counts <- sapply(subpop[, c("genre", "gender", "race")],
  function(x) length(unique(na.omit(x))))

if(any(factor_counts < 2)) {
  message(paste("Skipping", decade_label, "due to single-level factor:",
    names(factor_counts)[factor_counts < 2]))
  next
}

first_model <- lm(data = subpop, rating ~ age + race + gender + genre + energy
  + danceability + happiness + bpm + loudness_db + length_sec)

second_model <- lm(data = subpop, rating ~ age + gender + race + age*genre
  + age*energy + age*danceability + age*happiness + age*bpm
  + age*loudness_db + age*length_sec + race*genre + race*energy
  + race*danceability + race*happiness + race*bpm + race*loudness_db
  + race*length_sec + gender*genre + gender*energy + gender*danceability
  + gender*happiness + gender*bpm + gender*loudness_db + gender*length_sec)

if(anova(first_model, second_model)$"Pr(>F)"[2] <= 0.1){
  optimal_model <- step(second_model, direction="backward")
} else {
  optimal_model <- step(first_model, direction="backward")
}

model_results[[decade_label]] <- list(
  "first" = first_model,
  "second" = second_model,
  "optimal" = optimal_model
)

print(decade_label)
summary(first_model)
summary(second_model)
summary(optimal_model)
}


```

```
## Skipping 1950s - Not enough data points.
```

```
## Skipping 2020s due to single-level factor: genre
```

```
library("GGally")
```

```
ggpairs(data = cleaned_data[,c("energy", "danceability", "happiness", "loudness_db", "bpm", "length_sec")])
```



Loop through each decade in the model_results list

```
for (decade in names(model_results)) {
```

Check if the 'optimal' model exists for this decade

```
if (!is.null(model_results[[decade]]$optimal)) {
```

```
  cat("\n", "=====", "\n")
```

```
  cat(" SUMMARY FOR OPTIMAL MODEL:", decade, "\n")
```

```
  cat("=====", "\n")
```

Print the summary

```
  print(summary(model_results[[decade]]$optimal))
```

```
} else {
```

```
  cat("\nNo optimal model found for:", decade, "\n")
```

```
}
```

```
}
```

```
##
```

```
## =====
```

```
## SUMMARY FOR OPTIMAL MODEL: 1960s
```

```
## =====
```

```
##
```

```
## Call:
```

```
## lm(formula = rating ~ age + gender + race + genre + energy +
```

```
## danceability + happiness + length_sec + race:genre + race:energy +
```

```
## race:danceability + gender:danceability + gender:happiness,
```

```
## data = subpop)
```

```

##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.3064 -1.2960  0.1313  1.2436  4.2045
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    6.235892   2.001360   3.116  0.00214 **
## age           -0.036813   0.021122  -1.743  0.08309 .
## genderNot Male -1.820687   1.293523  -1.408  0.16101
## raceWhite      1.604179   1.795663   0.893  0.37287
## genreJazz      1.473039   1.332354   1.106  0.27040
## genreOther     1.978211   0.999147   1.980  0.04926 *
## genrePop       0.988531   0.889417   1.111  0.26788
## genreRock      0.260014   0.629236   0.413  0.67994
## energy        -0.005765   0.015336  -0.376  0.70743
## danceability   -0.013030   0.022465  -0.580  0.56264
## happiness      0.004845   0.009013   0.538  0.59154
## length_sec     0.008965   0.003791   2.365  0.01912 *
## raceWhite:genreJazz -1.353373  2.222521  -0.609  0.54334
## raceWhite:genreOther -3.249204  1.187123  -2.737  0.00683 **
## raceWhite:genrePop -1.786329  0.784348  -2.277  0.02395 *
## raceWhite:genreRock      NA         NA      NA      NA
## raceWhite:energy    0.024399  0.017347   1.407  0.16132
## raceWhite:danceability -0.046982  0.022464  -2.091  0.03791 *
## genderNot Male:danceability 0.076584  0.023552   3.252  0.00137 **
## genderNot Male:happiness -0.030694  0.016106  -1.906  0.05830 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.797 on 178 degrees of freedom
## Multiple R-squared:  0.314, Adjusted R-squared:  0.2446
## F-statistic: 4.527 on 18 and 178 DF, p-value: 4.551e-08
##
##
## =====
## SUMMARY FOR OPTIMAL MODEL: 1970s
## =====
##
## Call:
## lm(formula = rating ~ age + race + genre + length_sec, data = subpop)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.1615 -1.0209  0.0511  1.0964  5.0519
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    5.250133   0.992268   5.291 2.77e-07 ***
## age           -0.033100   0.020798  -1.591  0.112830
## raceWhite     -0.430183   0.299921  -1.434  0.152800
## genreJazz     -0.547224   1.924576  -0.284  0.776402
## genreOther     0.979427   0.660001   1.484  0.139144
## genrePop      -0.607790   0.641502  -0.947  0.344375

```

```

## genreRock    0.182210    0.658709    0.277 0.782315
## length_sec   0.009494    0.002515    3.775 0.000202 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.808 on 237 degrees of freedom
## Multiple R-squared:  0.2243, Adjusted R-squared:  0.2013
## F-statistic: 9.787 on 7 and 237 DF,  p-value: 1.045e-10
##
##
## =====
## SUMMARY FOR OPTIMAL MODEL: 1980s
## =====
##
## Call:
## lm(formula = rating ~ energy + danceability + happiness, data = subpop)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.2377 -1.0513 -0.0178  0.9701  4.2309
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.960976    0.549218   9.033  <2e-16 ***
## energy         0.009412    0.005806   1.621   0.1064
## danceability   0.015115    0.009255   1.633   0.1038
## happiness     -0.011560    0.005495  -2.104   0.0365 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.505 on 224 degrees of freedom
## Multiple R-squared:  0.02279, Adjusted R-squared:  0.009699
## F-statistic: 1.741 on 3 and 224 DF,  p-value: 0.1594
##
##
## =====
## SUMMARY FOR OPTIMAL MODEL: 1990s
## =====
##
## Call:
## lm(formula = rating ~ race + gender + energy + happiness + loudness_db,
##     data = subpop)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.2837 -0.9790 -0.0504  0.8700  4.4145
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.465767    0.942926  10.039  < 2e-16 ***
## raceWhite     -0.469197    0.277569  -1.690   0.09337 .
## genderNot Male  0.720941    0.253791   2.841   0.00523 **
## energy        -0.035118    0.010721  -3.276   0.00135 **
## happiness      0.016936    0.006634   2.553   0.01185 *

```

```

## loudness_db      0.311544    0.066407    4.691 6.82e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.469 on 129 degrees of freedom
## Multiple R-squared:  0.2445, Adjusted R-squared:  0.2152
## F-statistic: 8.349 on 5 and 129 DF,  p-value: 7.287e-07
##
##
## =====
## SUMMARY FOR OPTIMAL MODEL: 2000s
## =====
##
## Call:
## lm(formula = rating ~ age + gender + race + genre + energy +
##     danceability + happiness + bpm + length_sec + age:danceability +
##     race:genre + race:danceability + race:happiness + race:bpm +
##     race:length_sec + gender:danceability + gender:happiness +
##     gender:bpm, data = subpop)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.80175 -0.88847 -0.00231  1.02953  2.93881
##
## Coefficients: (2 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      14.254799    5.689019   2.506 0.013714 *
## age              -0.262392    0.164223  -1.598 0.113013
## genderNot Male    -0.551642    2.208098  -0.250 0.803196
## raceWhite        -13.479191    3.724145  -3.619 0.000451 ***
## genreOther         0.362481    1.668695   0.217 0.828443
## genrePop           0.218854    1.608297   0.136 0.892013
## genreRock          0.090807    1.604975   0.057 0.954986
## energy            0.018882    0.011341   1.665 0.098826 .
## danceability      -0.083852    0.065397  -1.282 0.202519
## happiness          0.001354    0.011393   0.119 0.905583
## bpm               -0.014549    0.008120  -1.792 0.075985 .
## length_sec        -0.009048    0.005130  -1.764 0.080619 .
## age:danceability   0.003737    0.002195   1.703 0.091492 .
## raceWhite:genreOther -1.186109    0.866090  -1.369 0.173685
## raceWhite:genrePop      NA         NA         NA      NA
## raceWhite:genreRock     NA         NA         NA      NA
## raceWhite:danceability  0.084078    0.031491   2.670 0.008759 **
## raceWhite:happiness    -0.025028    0.018899  -1.324 0.188199
## raceWhite:bpm           0.035295    0.012426   2.840 0.005385 **
## raceWhite:length_sec    0.019703    0.011520   1.710 0.090064 .
## genderNot Male:danceability -0.041970    0.024007  -1.748 0.083262 .
## genderNot Male:happiness  0.021226    0.014487   1.465 0.145771
## genderNot Male:bpm       0.019019    0.010619   1.791 0.076076 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.39 on 108 degrees of freedom
## Multiple R-squared:  0.3417, Adjusted R-squared:  0.2198

```

```
## F-statistic: 2.803 on 20 and 108 DF, p-value: 0.0003367
##
##
## =====
## SUMMARY FOR OPTIMAL MODEL: 2010s
## =====
##
## Call:
## lm(formula = rating ~ gender + energy + happiness + length_sec,
## data = subpop)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.2684 -1.0215  0.1378  0.9163  3.4440
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    4.788777    1.188134   4.031 0.000103 ***
## genderNot Male    0.459667    0.298784   1.538 0.126808
## energy         -0.024064    0.009966  -2.415 0.017405 *
## happiness        0.012193    0.007833   1.557 0.122440
## length_sec      0.009668    0.004602   2.101 0.037952 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.594 on 110 degrees of freedom
## Multiple R-squared:  0.09648, Adjusted R-squared:  0.06363
## F-statistic: 2.937 on 4 and 110 DF, p-value: 0.02378
```

Preliminary Exploration of Data through Visual Analysis

Plots of Individual Relationships

The plots show the relationship between the predictors and the response variable, rating.

```
library(ggplot2)
library(patchwork)

##
## Attaching package: 'patchwork'
##
## The following object is masked from 'package:MASS':
##
##      area

plot_age <- ggplot(cleaned_data, aes(x = age, y = rating)) +
  geom_point(color = "black", size = 1.5) +
  geom_smooth(method = "lm", color = "red", se = F) +
  theme_classic() +
  ggtitle("Rating vs Age")

plot_energy <- ggplot(cleaned_data, aes(x = energy, y = rating)) +
  geom_point(color = "black", size = 1.5) +
  geom_smooth(method = "lm", color = "red", se = F) +
  theme_classic() +
  ggtitle("Rating vs Energy")
```

```
plot_danceability <- ggplot(cleaned_data, aes(x = danceability, y = rating)) +
  geom_point(color = "black", size = 1.5) +
  geom_smooth(method = "lm", color = "red", se = F) +
  theme_classic() +
  ggtitle("Rating vs Danceability")
```

```
plot_happiness <- ggplot(cleaned_data, aes(x = happiness, y = rating)) +
  geom_point(color = "black", size = 1.5) +
  geom_smooth(method = "lm", color = "red", se = F) +
  theme_classic() +
  ggtitle("Rating vs Happiness")
```


```
plot_age + plot_energy + plot_danceability + plot_happiness
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

```
## `geom_smooth()` using formula = 'y ~ x'
```

```
## `geom_smooth()` using formula = 'y ~ x'
```



```
plot_bpm <- ggplot(cleaned_data, aes(x = bpm, y = rating)) +
  geom_point(color = "black", size = 1.5) +
  geom_smooth(method = "lm", color = "red", se = F) +
  theme_classic() +
  ggtitle("Rating vs BPM")
```

```
plot_loudness_db <- ggplot(cleaned_data, aes(x = loudness_db, y = rating)) +
  geom_point(color = "black", size = 1.5) +
```

```


geom_smooth(method = "lm", color = "red", se = F) +
theme_classic() +
ggtitle("Rating vs Loudness(db)")

plot_length_sec <- ggplot(cleaned_data, aes(x = length_sec, y = rating)) +
  geom_point(color = "black", size = 1.5) +
  geom_smooth(method = "lm", color = "red", se = F) +
  theme_classic() +
  ggtitle("Rating vs Song Length(sec)")

plot_bpm + plot_loudness_db + plot_length_sec

## `geom_smooth()` using formula = 'y ~ x'
## `geom_smooth()` using formula = 'y ~ x'
## `geom_smooth()` using formula = 'y ~ x'

```



```


plot_race <- ggplot(cleaned_data, aes(x = race, y = rating, fill = race)) +
  geom_boxplot() +
  theme_classic() +
  ggtitle("Rating vs Race") +
  guides(fill = "none")

plot_gender <- ggplot(cleaned_data, aes(x = gender, y = rating, fill = gender)) +
  geom_boxplot() +
  theme_classic() +
  ggtitle("Rating vs Gender") +
  guides(fill = "none")

```


```
plot_genre <- ggplot(cleaned_data, aes(x = genre, y = rating, fill = genre)) +
  geom_boxplot() +
  theme_classic() +
  ggtitle("Rating vs Genre")
```

```
plot_race + plot_gender + plot_genre
```



As expected, banding is observed in the scatterplots because the response variable, rating, is discrete. Since race, gender, and genre are categorical predictors, the relationship with rating is displayed by boxplots instead.

```
ggplot(cleaned_data, aes(x = rating)) +
  geom_histogram(
    binwidth = 0.5,
    boundary = 0.5,
    fill = "black",
    color = "grey",
    linewidth = 0.3
  ) +
  labs(
    title = "Histogram of Song Ratings",
    x = "Rating",
    y = "Frequency"
  ) +
  theme_classic()
```



However, the histogram of ratings appears approximately symmetric and unimodal, resembling a normal distribution. While the ratings are discrete (integer-valued) and there are many categorical predictors, the distribution spans multiple levels, creating a shape that approximates continuity. This supports the treating of rating as a continuous response variable in the regression analysis.

Plots of Grouped Relationships

Data visualizations for grouped relationships help us examine whether relationships between predictors and ratings differ across categories. These plots justify the inclusion of categorical predictors in the regression models and help identify potential interaction effects.

```
model_60s <- model_results[["1960s"]]$optimal
model_70s <- model_results[["1970s"]]$optimal
model_80s <- model_results[["1980s"]]$optimal
model_90s <- model_results[["1990s"]]$optimal
model_2000s <- model_results[["2000s"]]$optimal
model_2010s <- model_results[["2010s"]]$optimal
```

```
summary(model_60s)$r.squared
```

```
## [1] 0.3140162
```

```
summary(model_60s)$adj.r.squared
```

```
## [1] 0.244647
```

```
summary(model_70s)$r.squared
```

```
## [1] 0.2242527
```

```

summary(model_70s)$adj.r.squared

## [1] 0.2013403
summary(model_80s)$r.squared

## [1] 0.02278648
summary(model_80s)$adj.r.squared

## [1] 0.009698796
summary(model_90s)$r.squared

## [1] 0.2444873
summary(model_90s)$adj.r.squared

## [1] 0.2152039
summary(model_2000s)$r.squared

## [1] 0.3416655
summary(model_2000s)$adj.r.squared

## [1] 0.2197517
summary(model_2010s)$r.squared

## [1] 0.09648348
summary(model_2010s)$adj.r.squared

## [1] 0.06362834
vars_by_decade <- lapply(model_results, function(x) {
  attr(terms(x$optimal), "term.labels")
})

all_vars <- unique(unlist(vars_by_decade))

presence_matrix <- sapply(vars_by_decade, function(vars) {
  all_vars %in% vars
})

rownames(presence_matrix) <- all_vars
presence_matrix

##           1960s 1970s 1980s 1990s 2000s 2010s
## age           TRUE  TRUE FALSE FALSE  TRUE FALSE
## gender        TRUE FALSE FALSE  TRUE  TRUE  TRUE
## race          TRUE  TRUE FALSE  TRUE  TRUE FALSE
## genre         TRUE  TRUE FALSE FALSE  TRUE FALSE
## energy        TRUE FALSE  TRUE  TRUE  TRUE  TRUE
## danceability  TRUE FALSE  TRUE FALSE  TRUE FALSE
## happiness     TRUE FALSE  TRUE  TRUE  TRUE  TRUE
## length_sec    TRUE  TRUE FALSE FALSE  TRUE  TRUE
## race:genre    TRUE FALSE FALSE FALSE  TRUE FALSE
## race:energy   TRUE FALSE FALSE FALSE FALSE FALSE

```

```
## race:danceability    TRUE FALSE FALSE FALSE  TRUE FALSE
## gender:danceability  TRUE FALSE FALSE FALSE  TRUE FALSE
## gender:happiness    TRUE FALSE FALSE FALSE  TRUE FALSE
## loudness_db         FALSE FALSE FALSE  TRUE FALSE FALSE
## bpm                 FALSE FALSE FALSE FALSE  TRUE FALSE
## age:danceability    FALSE FALSE FALSE FALSE  TRUE FALSE
## race:happiness      FALSE FALSE FALSE FALSE  TRUE FALSE
## race:bpm            FALSE FALSE FALSE FALSE  TRUE FALSE
## race:length_sec     FALSE FALSE FALSE FALSE  TRUE FALSE
## gender:bpm          FALSE FALSE FALSE FALSE  TRUE FALSE
```

```
rownames(presence_matrix)[rowSums(presence_matrix) == ncol(presence_matrix)]
```

```
## character(0)
```

```
presence_df <- as.data.frame(as.table(presence_matrix))
```

```
ggplot(presence_df, aes(x = Var2, y = Var1, fill = Freq)) +
  geom_tile() +
  scale_fill_manual(values = c("black", "#23DC6D")) +
  labs(x = "Decade", y = "Variable", fill = "Included") +
  theme_minimal()
```



```
true_counts <- rowSums(presence_matrix)
true_counts
```

```
##          age          gender          race          genre
##          3           4           4           3
##    energy  danceability  happiness  length_sec
```

```
##           5           3           5           4
##      race:genre      race:energy  race:danceability  gender:danceability
##           2           1           2           2
##      gender:happiness      loudness_db      bpm      age:danceability
##           2           1           1           1
##      race:happiness      race:bpm      race:length_sec      gender:bpm
##           1           1           1           1
```

```
rownames(presence_matrix)[true_counts == max(true_counts)]
```

```
## [1] "energy"      "happiness"
```

```
false_counts <- rowSums(!presence_matrix)
false_counts
```

```
##           age           gender           race           genre
##           3           2           2           3
##      energy      danceability      happiness      length_sec
##           1           3           1           2
##      race:genre      race:energy  race:danceability  gender:danceability
##           4           5           4           4
##      gender:happiness      loudness_db      bpm      age:danceability
##           4           5           5           5
##      race:happiness      race:bpm      race:length_sec      gender:bpm
##           5           5           5           5
```

```
rownames(presence_matrix)[false_counts == max(false_counts)]
```

```
## [1] "race:energy"      "loudness_db"      "bpm"      "age:danceability"
## [5] "race:happiness"  "race:bpm"      "race:length_sec"  "gender:bpm"
```

```
summary_table <- data.frame(
  TRUE_count = rowSums(presence_matrix),
  FALSE_count = rowSums(!presence_matrix)
)
```


```
summary_table[order(-summary_table$TRUE_count), ]
```

```
##           TRUE_count  FALSE_count
## energy              5             1
## happiness           5             1
## gender              4             2
## race                4             2
## length_sec          4             2
## age                 3             3
## genre               3             3
## danceability         3             3
## race:genre           2             4
## race:danceability     2             4
## gender:danceability    2             4
## gender:happiness      2             4
## race:energy           1             5
## loudness_db           1             5
## bpm                  1             5
## age:danceability      1             5
## race:happiness        1             5
## race:bpm              1             5
```

```
## race:length_sec      1      5
## gender:bpm           1      5
```



```
par(mfrow = c(1,2))
plot(model_60s, which = 1:2, main = "1960s")
```

```
## Warning: not plotting observations with leverage one:
##      48
```





```
plot(model_70s, which = 1:2, main = "1970s")
```



```
## Warning: not plotting observations with leverage one:
##      239
```



```
plot(model_80s, which = 1:2, main = "1980s")
```





```
plot(model_90s, which = 1:2, main = "1990s")
```




```
plot(model_2000s, which = 1:2, main = "2000s")
```

```
## Warning: not plotting observations with leverage one:
## 4
```



```
plot(model_2010s, which = 1:2, main = "2010s")
```



The fitted regression lines indicate that the relationship between song length and rating varies by genre. In particular, certain genres (e.g., Other and Country) exhibit steeper positive slopes, while others (e.g., Rock and Pop) show weaker or slightly negative trends, suggesting a potential interaction between length and genre.

```
library(broom)
```

```
## Warning: package 'broom' was built under R version 4.4.3
```

```
coef_1960s <- tidy(model_60s, conf.int = TRUE) %>% mutate(decade = "1960s")
coef_1970s <- tidy(model_70s, conf.int = TRUE) %>% mutate(decade = "1970s")
coef_1980s <- tidy(model_80s, conf.int = TRUE) %>% mutate(decade = "1980s")
coef_1990s <- tidy(model_90s, conf.int = TRUE) %>% mutate(decade = "1990s")
coef_2000s <- tidy(model_2000s, conf.int = TRUE) %>% mutate(decade = "2000s")
coef_2010s <- tidy(model_2010s, conf.int = TRUE) %>% mutate(decade = "2010s")

#Appending Data not Represented
#1960s
coef_1960s <- rbind(coef_1960s, list("genreCountry", 0, 0, 0, 0, 0, 0, "1960s"))

#1970s
coef_1970s <- rbind(coef_1970s, list("genreCountry", 0, 0, 0, 0, 0, 0, "1970s"))

#1980s
coef_1980s <- rbind(coef_1980s, list("genreCountry", 0, 0, 0, 0, 0, 0, "1980s"))
coef_1980s <- rbind(coef_1980s, list("genreOther", 0, 0, 0, 0, 0, 0, "1980s"))
coef_1980s <- rbind(coef_1980s, list("genreJazz", 0, 0, 0, 0, 0, 0, "1980s"))
coef_1980s <- rbind(coef_1980s, list("genreRock", 0, 0, 0, 0, 0, 0, "1980s"))
coef_1980s <- rbind(coef_1980s, list("genrePop", 0, 0, 0, 0, 0, 0, "1980s"))
```

```

#1990s
coef_1990s <- rbind(coef_1990s, list("genreCountry", 0, 0, 0, 0, 0, 0, "1990s"))
coef_1990s <- rbind(coef_1990s, list("genreOther", 0, 0, 0, 0, 0, 0, "1990s"))
coef_1990s <- rbind(coef_1990s, list("genreJazz", 0, 0, 0, 0, 0, 0, "1990s"))
coef_1990s <- rbind(coef_1990s, list("genreRock", 0, 0, 0, 0, 0, 0, "1990s"))
coef_1990s <- rbind(coef_1990s, list("genrePop", 0, 0, 0, 0, 0, 0, "1990s"))

#2000s
coef_2000s <- rbind(coef_2000s, list("genreCountry", 0, 0, 0, 0, 0, 0, "2000s"))
coef_2000s <- rbind(coef_2000s, list("genreJazz", 0, 0, 0, 0, 0, 0, "2000s"))

#2010s
coef_2010s <- rbind(coef_2010s, list("genreCountry", 0, 0, 0, 0, 0, 0, "2010s"))
coef_2010s <- rbind(coef_2010s, list("genreOther", 0, 0, 0, 0, 0, 0, "2010s"))
coef_2010s <- rbind(coef_2010s, list("genreJazz", 0, 0, 0, 0, 0, 0, "2010s"))
coef_2010s <- rbind(coef_2010s, list("genreRock", 0, 0, 0, 0, 0, 0, "2010s"))
coef_2010s <- rbind(coef_2010s, list("genrePop", 0, 0, 0, 0, 0, 0, "2010s"))


coef_data <- bind_rows(
  coef_1960s,
  coef_1970s,
  coef_1980s,
  coef_1990s,
  coef_2000s,
  coef_2010s
)

coef_data <- coef_data %>%
  filter(term != "(Intercept)") %>%
  filter(grepl("^genre", term)) %>%
  mutate(term = gsub("^genre", "", term))

coef_data$term <- factor(
  coef_data$term,
  levels = c("Jazz", "Pop", "Rock", "Country", "HipHop", "Other")
)

ggplot(coef_data, aes(x = estimate, y = term, color = term)) +
  geom_vline(xintercept = 0, linetype = "dashed") +
  geom_pointrange(aes(xmin = conf.low, xmax = conf.high), size = 0.8) +
  facet_wrap(~ decade) +
  theme_classic() +
  labs(
    title = "Genre Effects on Song Ratings by Decade",
    x = "Coefficient Estimate",
    y = "Genre"
  )

```



Country will always be 0 since it is our reference level.

model_60s

```
##
## Call:
## lm(formula = rating ~ age + gender + race + genre + energy +
##     danceability + happiness + length_sec + race:genre + race:energy +
##     race:danceability + gender:danceability + gender:happiness,
##     data = subpop)
##
## Coefficients:
##             (Intercept)                age
##             6.235892                -0.036813
##      genderNot Male          raceWhite
##      -1.820687                1.604179
##      genreJazz              genreOther
##      1.473039                1.978211
##      genrePop              genreRock
##      0.988531                0.260014
##      energy              danceability
##      -0.005765              -0.013030
##      happiness            length_sec
##      0.004845                0.008965
##      raceWhite:genreJazz    raceWhite:genreOther
##      -1.353373              -3.249204
##      raceWhite:genrePop     raceWhite:genreRock
##      -1.786329                NA
```

```

##           raceWhite:energy      raceWhite:danceability
##                0.024399                -0.046982
## genderNot Male:danceability  genderNot Male:happiness
##                0.076584                -0.030694
model_70s

##
## Call:
## lm(formula = rating ~ age + race + genre + length_sec, data = subpop)
##
## Coefficients:
## (Intercept)      age      raceWhite      genreJazz      genreOther      genrePop
##    5.250133   -0.033100   -0.430183   -0.547224     0.979427   -0.607790
##  genreRock    length_sec
##    0.182210     0.009494
model_80s

##
## Call:
## lm(formula = rating ~ energy + danceability + happiness, data = subpop)
##
## Coefficients:
## (Intercept)      energy  danceability      happiness
##    4.960976     0.009412     0.015115   -0.011560
model_90s

##
## Call:
## lm(formula = rating ~ race + gender + energy + happiness + loudness_db,
##     data = subpop)
##
## Coefficients:
## (Intercept)      raceWhite  genderNot Male      energy      happiness
##    9.465777     -0.46920      0.72094     -0.03512     0.01694
##  loudness_db
##    0.31154
model_2000s

##
## Call:
## lm(formula = rating ~ age + gender + race + genre + energy +
##     danceability + happiness + bpm + length_sec + age:danceability +
##     race:genre + race:danceability + race:happiness + race:bpm +
##     race:length_sec + gender:danceability + gender:happiness +
##     gender:bpm, data = subpop)
##
## Coefficients:
## (Intercept)      age
##    14.254799     -0.262392
##  genderNot Male      raceWhite
##   -0.551642     -13.479191
##    genreOther      genrePop
##     0.362481      0.218854

```

```
##          genreRock          energy
##          0.090807          0.018882
##          danceability      happiness
##          -0.083852          0.001354
##          bpm              length_sec
##          -0.014549          -0.009048
##          age:danceability    raceWhite:genreOther
##          0.003737            -1.186109
##          raceWhite:genrePop    raceWhite:genreRock
##          NA                  NA
##          raceWhite:danceability  raceWhite:happiness
##          0.084078            -0.025028
##          raceWhite:bpm          raceWhite:length_sec
##          0.035295            0.019703
## genderNot Male:danceability  genderNot Male:happiness
##          -0.041970            0.021226
##          genderNot Male:bpm
##          0.019019
```

```
model_2010s
```

```
##
## Call:
## lm(formula = rating ~ gender + energy + happiness + length_sec,
##     data = subpop)
##
## Coefficients:
## (Intercept)  genderNot Male      energy      happiness      length_sec
##    4.788777      0.459667    -0.024064     0.012193     0.009668
```

```
anova(model_results[["1960s"]] $\$$ first, model_results[["1960s"]] $\$$ second)
```

```
## Analysis of Variance Table
##
## Model 1: rating ~ age + race + gender + genre + energy + danceability +
## happiness + bpm + loudness_db + length_sec
## Model 2: rating ~ age + gender + race + age * genre + age * energy + age *
## danceability + age * happiness + age * bpm + age * loudness_db +
## age * length_sec + race * genre + race * energy + race *
## danceability + race * happiness + race * bpm + race * loudness_db +
## race * length_sec + gender * genre + gender * energy + gender *
## danceability + gender * happiness + gender * bpm + gender *
## loudness_db + gender * length_sec
## Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      183 673.25
## 2      155 535.98 28    137.27 1.4178 0.09454 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(model_results[["1970s"]] $\$$ first, model_results[["1970s"]] $\$$ second)
```

```
## Analysis of Variance Table
##
## Model 1: rating ~ age + race + gender + genre + energy + danceability +
## happiness + bpm + loudness_db + length_sec
## Model 2: rating ~ age + gender + race + age * genre + age * energy + age *
```

```

##      danceability + age * happiness + age * bpm + age * loudness_db +
##      age * length_sec + race * genre + race * energy + race *
##      danceability + race * happiness + race * bpm + race * loudness_db +
##      race * length_sec + gender * genre + gender * energy + gender *
##      danceability + gender * happiness + gender * bpm + gender *
##      loudness_db + gender * length_sec
##      Res.Df      RSS Df Sum of Sq      F Pr(>F)
## 1      231 764.63
## 2      204 666.81 27      97.812 1.1083 0.3327
anova(model_results[["1980s"]] $\$$ first, model_results[["1980s"]] $\$$ second)

## Analysis of Variance Table
##
## Model 1: rating ~ age + race + gender + genre + energy + danceability +
##      happiness + bpm + loudness_db + length_sec
## Model 2: rating ~ age + gender + race + age * genre + age * energy + age *
##      danceability + age * happiness + age * bpm + age * loudness_db +
##      age * length_sec + race * genre + race * energy + race *
##      danceability + race * happiness + race * bpm + race * loudness_db +
##      race * length_sec + gender * genre + gender * energy + gender *
##      danceability + gender * happiness + gender * bpm + gender *
##      loudness_db + gender * length_sec
##      Res.Df      RSS Df Sum of Sq      F Pr(>F)
## 1      216 494.88
## 2      192 459.91 24      34.968 0.6083 0.9247
anova(model_results[["1990s"]] $\$$ first, model_results[["1990s"]] $\$$ second)

## Analysis of Variance Table
##
## Model 1: rating ~ age + race + gender + genre + energy + danceability +
##      happiness + bpm + loudness_db + length_sec
## Model 2: rating ~ age + gender + race + age * genre + age * energy + age *
##      danceability + age * happiness + age * bpm + age * loudness_db +
##      age * length_sec + race * genre + race * energy + race *
##      danceability + race * happiness + race * bpm + race * loudness_db +
##      race * length_sec + gender * genre + gender * energy + gender *
##      danceability + gender * happiness + gender * bpm + gender *
##      loudness_db + gender * length_sec
##      Res.Df      RSS Df Sum of Sq      F Pr(>F)
## 1      122 269.95
## 2       98 217.50 24      52.448 0.9846 0.493
anova(model_results[["2000s"]] $\$$ first, model_2000s)

## Analysis of Variance Table
##
## Model 1: rating ~ age + race + gender + genre + energy + danceability +
##      happiness + bpm + loudness_db + length_sec
## Model 2: rating ~ age + gender + race + genre + energy + danceability +
##      happiness + bpm + length_sec + age:danceability + race:genre +
##      race:danceability + race:happiness + race:bpm + race:length_sec +
##      gender:danceability + gender:happiness + gender:bpm
##      Res.Df      RSS Df Sum of Sq      F      Pr(>F)
## 1      116 270.33

```

```
## 2    108 208.61  8    61.719 3.994 0.0003515 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

anova(model_2000s, model_results[["2000s"]])$second)

## Analysis of Variance Table
##
## Model 1: rating ~ age + gender + race + genre + energy + danceability +
##   happiness + bpm + length_sec + age:danceability + race:genre +
##   race:danceability + race:happiness + race:bpm + race:length_sec +
##   gender:danceability + gender:happiness + gender:bpm
## Model 2: rating ~ age + gender + race + age * genre + age * energy + age *
##   danceability + age * happiness + age * bpm + age * loudness_db +
##   age * length_sec + race * genre + race * energy + race *
##   danceability + race * happiness + race * bpm + race * loudness_db +
##   race * length_sec + gender * genre + gender * energy + gender *
##   danceability + gender * happiness + gender * bpm + gender *
##   loudness_db + gender * length_sec
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      108 208.61
## 2       93 194.20 15    14.415 0.4602 0.9544

anova(model_results[["2010s"]])$first, model_2010s)

## Analysis of Variance Table
##
## Model 1: rating ~ age + race + gender + genre + energy + danceability +
##   happiness + bpm + loudness_db + length_sec
## Model 2: rating ~ gender + energy + happiness + length_sec
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1       103 274.94
## 2       110 279.47 -7    -4.5351 0.2427 0.9735

anova(model_2010s, model_results[["2010s"]])$second)

## Analysis of Variance Table
##
## Model 1: rating ~ gender + energy + happiness + length_sec
## Model 2: rating ~ age + gender + race + age * genre + age * energy + age *
##   danceability + age * happiness + age * bpm + age * loudness_db +
##   age * length_sec + race * genre + race * energy + race *
##   danceability + race * happiness + race * bpm + race * loudness_db +
##   race * length_sec + gender * genre + gender * energy + gender *
##   danceability + gender * happiness + gender * bpm + gender *
##   loudness_db + gender * length_sec
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1       110 279.47
## 2        81 206.34 29    73.132 0.9899 0.4938
```